

Data Structure Programming Project #5

郭建志

Image

- You want to count elements
- You don't need exact results

Problem

- Given:
- keys in many documents
- Goal:
- Count the key frequency
- Bounded error is allowed
- Constraint:
- Limited storage and limited computation

Simple Solution

- Construct a map from elements to counts
- Balanced binary tree:
`map` in C++
- Hash table:
`unordered_map` in C++
- You may want to use libraries Guava or FastUtil in Java for convenience and better performance

Problem

- The number of distinct elements might be very large
- You have a limited memory space
- For real-time applications you need runtime guarantees

Solution: Count-Min Sketch

- The trick: don't store the distinct elements, but just the counters
- Create an integer array of length x initially filled with 0s
- Each incoming element gets mapped to a number between 0 and x
- The corresponding counter in the array gets incremented
- To query an element's count, simply return the integer value at its position


Solution: Count-Min Sketch

- The trick: don't store the distinct elements, but just the counters
- Create an integer array of length x initially filled with 0s
- Each incoming element gets mapped to a number
- The counter at that index is incremented
- To query an element's count, simply return the integer value at its position

You are completely right:
There will be collisions!

Solution: Count-Min Sketch

- Use **multiple** arrays with different hash functions to compute the index
- When queried, return the minimum of the numbers the array



Solution: Count-Min Sketch

- Use **multiple** arrays with different hash functions to compute the index
- When queried, return the minimum of the numbers the array

You are completely right:
There will still be collisions!

Solution: Count-Min Sketch

- Use **multiple** arrays with different hash functions to compute the index
- When queried, return the minimum of the numbers the array


You are completely right:
There will still be collisions!
... but less

Some properties

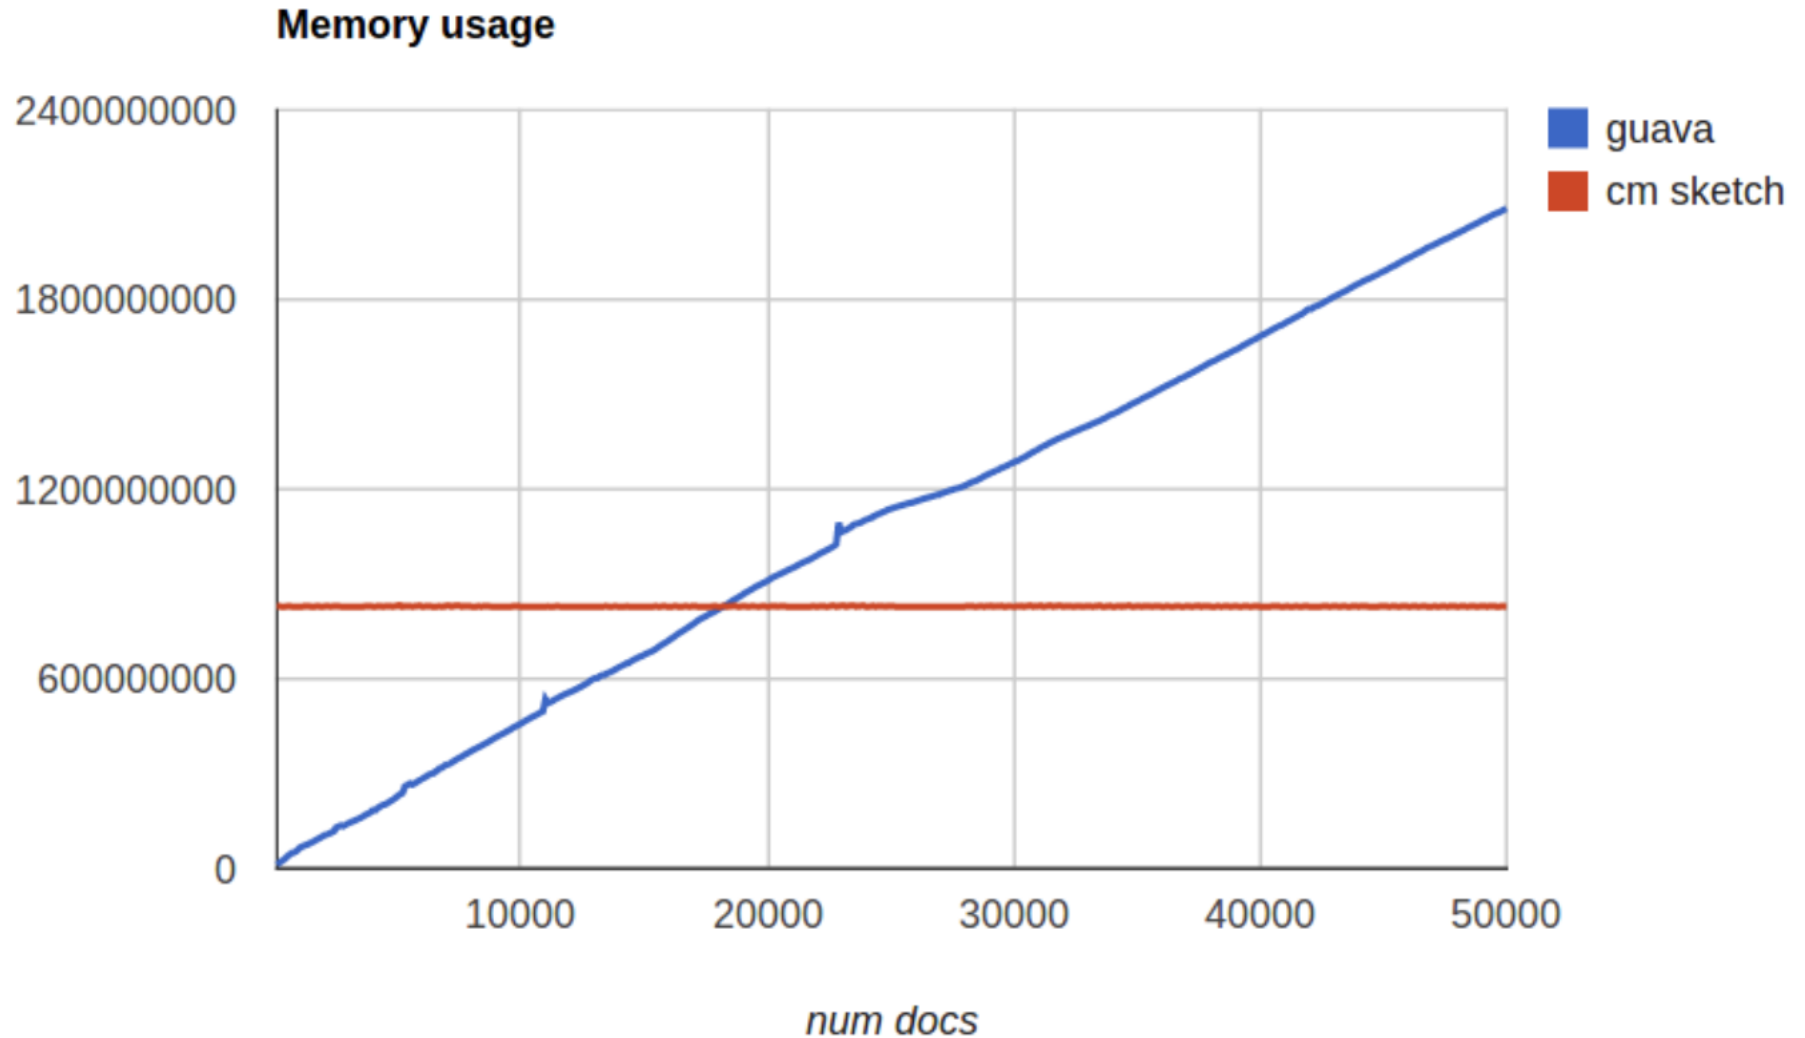
- Only over-estimates, never under-estimates the true count
- Has a **constant** memory and time consumption independent of the number of elements
- The relative error may be high for low-frequent elements

To be practical: Conservative Update

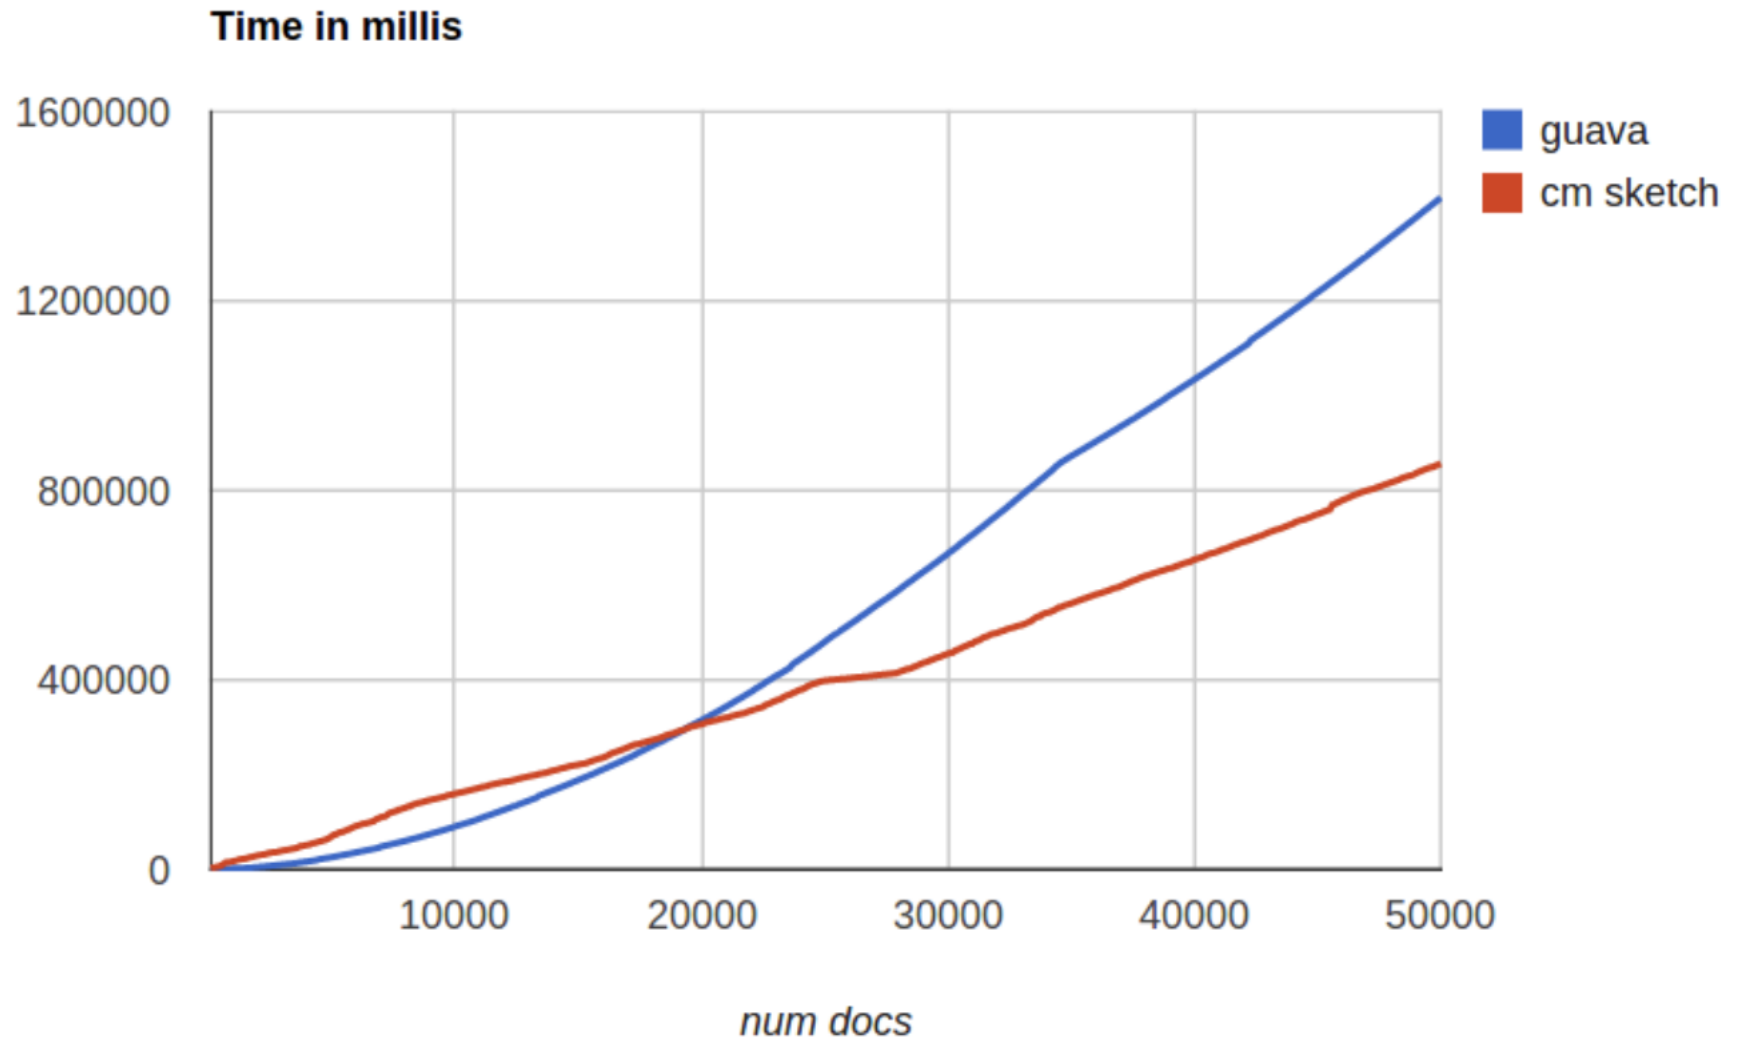
- Instead of incrementing every counter, **only increment the ones which need an update**
- Experiments show a strong error reduction using this heuristic



Comparison



Comparison



Usage

- General:
 - Every counting problem on large data sets, where errors are acceptable
 - Counting elements in data streams (e.g., trend detection on Twitter, Facebook, and news)

Usage

- Mine:
 - Feature selection for large sets using pointwise mutual information (PMI)
 - PMI has been used for finding associations between words
 - All kinds of feature statistics like word co-occurrences

You need to implement:

```
void init(int **map, int r, int c, int *a, int *b, int p)
```

1. Create a 2D array with r rows and c columns for pointer map
2. Create an array with r elements uniformly chosen from [1, p-1] for pointer a
3. Create an array with r elements uniformly chosen from [1, p-1] for pointer b (note: a[i] and b[i] should be independent)

```
int myhash(char *str, int count, int r, int c, int p, int *a, int *b)
```

Use `hash` in `<string>` to covert str to an integer key

```
// note that 0 <= count <= r-1
```

```
return (a[count] * key + b[count]) % p % c;
```

```
void insert(int **map, int r, int c, int p, char *str, int *a, int *b)
```

1. Find the smallest `ones` in the following positions,
map[count][myhash(str, count, r, c, a, b)] for $0 \leq \text{count} \leq r-1$
2. Increment the smallest `ones` by 1

```
void query(int **map, int r, int c, int p, char *str, int *a, int *b)
```

1. Find the smallest `one` in the following positions,
map[count][myhash(str, count, r, c, a, b)] for $0 \leq \text{count} \leq r-1$
2. Return the smallest `one`

Input Sample: input.txt

```
10 10 1019
```

```
Data structures  
serve as the  
basis
```

```
for abstract  
data type. The  
abstract data  
type defines  
the logical  
form of the  
data type. The  
data structure  
implements the  
physical form  
of the data  
type.
```

```
#rows #cols prime
```

```
Text...
```

```
Note: your code  
should understand  
that "Data" is  
equivalent to "data"  
and ignore  
punctuation marks  
(e.g., commas)
```

You don't need to output anything, but..

Command:

data

6

serve

1

type

4

Input a key

Return a value

Note: The value is **allowed** to have a small error to some extent, since you are using cm sketch instead of binary search tree

Note

- **Deadline:**
1/3 Thu (有問題可以再調整)
- This project accounts for 15% of the total score
- You are not allowed to use "class" in STL to count the words
- You must implement a 2D array with the given size (i.e., #rows and #cols in input.txt) to count the words
- **E-course**
- **C or C++ Source code**
- You can read this paper if you are interested in the research field:
- <http://dimacs.rutgers.edu/~graham/pubs/papers/cmsoft.pdf>